

# prediction\_random\_forest

September 24, 2019

```
In [1]: import pandas as pd
```

```
# some lines have too many fields (?), so skip bad lines
imgatt = pd.read_csv("CUB_200_2011/attributes/image_attribute_labels.txt", sep='\s+',
                    header=None, error_bad_lines=False,
                    warn_bad_lines=False, usecols=[0,1,2],
                    names=['imgid', 'attid', 'present'])

# description from dataset README:
# The set of attribute labels as perceived by MTurkers for each image is contained
# in the file attributes/image_attribute_labels.txt,
# with each line corresponding to one image/attribute/worker triplet:

# <image_id> <attribute_id> <is_present> <certainty_id> <time>

# where <image_id>, <attribute_id>, <certainty_id> correspond to the IDs in images.txt
# attributes.txt, and attributes/certainties.txt
# respectively. <is_present> is 0 or 1 (1 denotes that the attribute is present).
# <time> denotes the time spent by the MTurker in seconds.
```

```
In [2]: imgatt.head()
```

```
Out[2]:
```

|   | imgid | attid | present |
|---|-------|-------|---------|
| 0 | 1     | 1     | 0       |
| 1 | 1     | 2     | 0       |
| 2 | 1     | 3     | 0       |
| 3 | 1     | 4     | 0       |
| 4 | 1     | 5     | 1       |

```
In [3]: imgatt.shape
```

```
Out[3]: (3677856, 3)
```

```
In [4]: # Image ID number 1 does not have attributes 1,2,3, or 4, but it does have attribute 5.
# The shape will tell us how many rows and columns we have.
# It has 3.7 million rows and three columns. This is not the actual formula that you want.
# You want attributes to be the columns, not rows.
# need to reorganize imgatt to have one row per imgid, and 312 columns (one column per attribute)
```

```

# with 1/0 in each cell representing
# if that imgid has that attribute or not
imgatt2 = imgatt.pivot(index='imgid', columns='attid', values='present')

```

In [5]: `imgatt2.head()`

```

Out[5]: attid  1    2    3    4    5    6    7    8    9   10   ...  303  304  305  \
imgid
1           0    0    0    0    1    0    0    0    0    0   ...    0    0    0
2           0    0    0    0    0    0    0    0    0    0   ...    0    0    0
3           0    0    0    0    1    0    0    0    0    0   ...    0    0    0
4           0    0    0    0    1    0    0    0    0    0   ...    0    0    0
5           0    0    0    0    1    0    0    0    0    0   ...    0    0    1

attid  306  307  308  309  310  311  312
imgid
1           0    0    1    0    0    0    0
2           0    0    0    0    0    0    0
3           0    0    1    0    0    1    0
4           1    0    0    1    0    0    0
5           0    0    0    0    0    0    0

[5 rows x 312 columns]

```

In [6]: `# feed data into a random forest. In the previous example, we have 312 columns and 312`  
`# which is ultimately`  
`# about 12,000 images or 12,000 different examples of birds:`  
`imgatt2.shape`

Out[6]: (11788, 312)

In [7]: `# need to load the answers, such as whether it;s a bird and which species it is.`  
`# Since it is an image class labels file, the separators`  
`# are spaces. There is no header row and the two columns are imgid and label.`  
`# We will be using set_index('imgid') to have the same result`  
`# produced by imgatt2.head(), where the rows are idenfitted by the image ID`

`# load the image true classes`

```

imglabels = pd.read_csv("CUB_200_2011/image_class_labels.txt", sep=' ', header=None, na
imglabels = imglabels.set_index('imgid')

```

`# decription from dataset README:`  
`# The ground truth class labels (bird species labels) for each image are contained`  
`# in the file image_class_labels.txt,`  
`# with each line corresponding to one image:`  
`# <image_id> <class_id>`  
`# where <image_id> and <class_id> correspond to the IDs in images.txt and classes.txt,`

In [8]: `imglabels.head()`

```
Out [8]:      label
imgid
1         1
2         1
3         1
4         1
5         1
```

```
In [9]: imglabels.shape
```

```
Out [9]: (11788, 1)
```

```
In [10]: df=imgatt2.join(imglabels)
df=df.sample(frac=1)
```

```
In [11]: df_att=df.iloc[:, :312]
df_label=df.iloc[:, 312:]
```

```
In [12]: df_att.head()
```

```
Out [12]:      1      2      3      4      5      6      7      8      9     10     ...    303    304    305  \
imgid
9426      0      0      0      0      0      0      1      0      0      0     ...      0      1      0
7056      0      0      0      0      0      0      0      1      0      0     ...      0      0      0
6321      0      0      0      0      0      0      0      0      0      0     ...      0      0      1
6359      0      0      0      0      0      0      1      0      0      0     ...      0      0      1
11409     0      1      0      0      0      0      0      0      0      0     ...      0      0      0
```

```
      306    307    308    309    310    311    312
imgid
9426      0      0      0      0      0      0      0
7056      0      0      0      0      1      0      0
6321      0      0      0      0      0      0      1
6359      0      0      0      0      0      0      1
11409     0      0      1      0      0      1      0
```

```
[5 rows x 312 columns]
```

```
In [62]: df_train_att = df_att[:8000]
df_train_label = df_label[:8000]
df_test_att = df_att[8000:]
df_test_label = df_label[8000:]
```

```
df_train_label = df_train_label['label']
df_test_label = df_test_label['label']
```

```
In [64]: from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier(max_features=50,random_state=0, n_estimators=100)
```

```
In [65]: clf.fit(df_train_att, df_train_label)
```

```
Out [65]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features=50, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=None,
                                oob_score=False, random_state=0, verbose=0, warm_start=False)
```

```
In [66]: print(clf.predict(df_train_att.head()))
```

```
[161 121 108 109 194]
```

```
In [67]: clf.score(df_test_att, df_test_label)
```

```
Out [67]: 0.4334741288278775
```

```
In [68]: from sklearn.metrics import confusion_matrix
pred_labels = clf.predict(df_test_att)
cm = confusion_matrix(df_test_label, pred_labels)
```

```
In [69]: cm
```

```
Out [69]: array([[ 6,  0,  1, ...,  0,  0,  0],
                 [ 1, 13,  0, ...,  0,  0,  0],
                 [ 3,  0,  4, ...,  0,  0,  0],
                 ...,
                 [ 0,  0,  1, ...,  4,  0,  0],
                 [ 0,  0,  0, ...,  0, 11,  0],
                 [ 0,  0,  0, ...,  0,  0, 14]], dtype=int64)
```

```
In [76]: import matplotlib.pyplot as plt
import itertools
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting normalize=True.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')
    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
```

```

tick_marks=np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=90)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max()/2
plt.tight_layout()
plt.ylabel("True label")
plt.xlabel("Predicted label")

```

```

In [77]: birds = pd.read_csv("CUB_200_2011/classes.txt",
                             sep='\s+', header=None, usecols=[1], names=['birdname'])
birds = birds['birdname']
birds

```

```

Out[77]: 0      001.Black_footed_Albatross
1      002.Laysan_Albatross
2      003.Sooty_Albatross
3      004.Groove_billed_Ani
4      005.Crested_Auklet
...
195     196.House_Wren
196     197.Marsh_Wren
197     198.Rock_Wren
198     199.Winter_Wren
199     200.Common_Yellowthroat
Name: birdname, Length: 200, dtype: object

```

```


In [78]: import numpy as np
np.set_printoptions(precision=2)
plt.figure(figsize=(60,60), dpi=300)
plot_confusion_matrix(cm, classes=birds, normalize=True)
plt.show()

```

```

Normalized confusion matrix
[[0.4  0.   0.07 ... 0.   0.   0. ]
 [0.04 0.5  0.   ... 0.   0.   0. ]
 [0.19 0.   0.25 ... 0.   0.   0. ]
 ...
 [0.   0.   0.06 ... 0.25 0.   0. ]
 [0.   0.   0.   ... 0.   0.5 0. ]
 [0.   0.   0.   ... 0.   0.  0.82]]

```



```
In [79]: from sklearn import tree
         clftree = tree.DecisionTreeClassifier()
         clftree.fit(df_train_att, df_train_label)
         clftree.score(df_test_att, df_test_label)
```

Out[79]: 0.26689545934530096

```
In [84]: from sklearn import svm
         clfsvm = svm.SVC(gamma='auto')
         clfsvm.fit(df_train_att, df_train_label)
         clfsvm.score(df_test_att, df_test_label)
```

Out[84]: 0.29197465681098206

```
In [86]: from sklearn.model_selection import cross_val_score
        scores = cross_val_score(clf, df_train_att, df_train_label, cv=5)
        print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std()*2))
```

Accuracy: 0.44 (+/- 0.01)

```
In [88]: scorestree = cross_val_score(clftree, df_train_att, df_train_label, cv=5)
        print("Accuracy: %0.2f (+/-%0.2f)" % (scorestree.mean(), scores.std()*2))
```

Accuracy: 0.25 (+/-0.01)

```
In [89]: scoressvm = cross_val_score(clfsvm, df_train_att, df_train_label, cv=5)
        print("Accuracy: %0.2f (+/-%0.2f)" % (scoressvm.mean(), scoressvm.std()*2))
```

Accuracy: 0.27 (+/-0.00)

```
In [91]: max_features_opts = range(5,50,5)
        n_estimators_opts = range(10,200,20)
        rf_params = np.empty((len(max_features_opts)*len(n_estimators_opts),4),float)
        i=0
        for max_features in max_features_opts:
            for n_estimators in n_estimators_opts:
                clf = RandomForestClassifier(max_features=max_features, n_estimators=n_estimators)
                scores = cross_val_score(clf, df_train_att, df_train_label, cv=5)
                rf_params[i,0] = max_features
                rf_params[i,1] = n_estimators
                rf_params[i,2] = scores.mean()
                rf_params[i,3] = scores.std()*2
                i+=1
            print("Max features: %d, num estimators: %d, accuracy: %0.2f (+/- %0.2f)" % \
                  (max_features, n_estimators, scores.mean(), scores.std()*2))
```

Max features: 5, num estimators: 10, accuracy: 0.27 (+/- 0.02)  
 Max features: 5, num estimators: 30, accuracy: 0.35 (+/- 0.02)  
 Max features: 5, num estimators: 50, accuracy: 0.39 (+/- 0.01)  
 Max features: 5, num estimators: 70, accuracy: 0.41 (+/- 0.01)  
 Max features: 5, num estimators: 90, accuracy: 0.42 (+/- 0.01)  
 Max features: 5, num estimators: 110, accuracy: 0.43 (+/- 0.01)  
 Max features: 5, num estimators: 130, accuracy: 0.44 (+/- 0.01)  
 Max features: 5, num estimators: 150, accuracy: 0.44 (+/- 0.01)  
 Max features: 5, num estimators: 170, accuracy: 0.44 (+/- 0.02)  
 Max features: 5, num estimators: 190, accuracy: 0.45 (+/- 0.01)  
 Max features: 10, num estimators: 10, accuracy: 0.29 (+/- 0.03)  
 Max features: 10, num estimators: 30, accuracy: 0.38 (+/- 0.02)  
 Max features: 10, num estimators: 50, accuracy: 0.41 (+/- 0.01)  
 Max features: 10, num estimators: 70, accuracy: 0.42 (+/- 0.01)

Max features: 10, num estimators: 90, accuracy: 0.43 (+/- 0.01)  
 Max features: 10, num estimators: 110, accuracy: 0.45 (+/- 0.02)  
 Max features: 10, num estimators: 130, accuracy: 0.44 (+/- 0.01)  
 Max features: 10, num estimators: 150, accuracy: 0.45 (+/- 0.01)  
 Max features: 10, num estimators: 170, accuracy: 0.45 (+/- 0.01)  
 Max features: 10, num estimators: 190, accuracy: 0.45 (+/- 0.01)  
 Max features: 15, num estimators: 10, accuracy: 0.31 (+/- 0.02)  
 Max features: 15, num estimators: 30, accuracy: 0.39 (+/- 0.02)  
 Max features: 15, num estimators: 50, accuracy: 0.42 (+/- 0.00)  
 Max features: 15, num estimators: 70, accuracy: 0.43 (+/- 0.02)  
 Max features: 15, num estimators: 90, accuracy: 0.44 (+/- 0.01)  
 Max features: 15, num estimators: 110, accuracy: 0.45 (+/- 0.01)  
 Max features: 15, num estimators: 130, accuracy: 0.45 (+/- 0.02)  
 Max features: 15, num estimators: 150, accuracy: 0.46 (+/- 0.01)  
 Max features: 15, num estimators: 170, accuracy: 0.46 (+/- 0.01)  
 Max features: 15, num estimators: 190, accuracy: 0.45 (+/- 0.02)  
 Max features: 20, num estimators: 10, accuracy: 0.32 (+/- 0.02)  
 Max features: 20, num estimators: 30, accuracy: 0.39 (+/- 0.02)  
 Max features: 20, num estimators: 50, accuracy: 0.42 (+/- 0.01)  
 Max features: 20, num estimators: 70, accuracy: 0.44 (+/- 0.01)  
 Max features: 20, num estimators: 90, accuracy: 0.44 (+/- 0.01)  
 Max features: 20, num estimators: 110, accuracy: 0.45 (+/- 0.01)  
 Max features: 20, num estimators: 130, accuracy: 0.45 (+/- 0.02)  
 Max features: 20, num estimators: 150, accuracy: 0.45 (+/- 0.02)  
 Max features: 20, num estimators: 170, accuracy: 0.45 (+/- 0.01)  
 Max features: 20, num estimators: 190, accuracy: 0.46 (+/- 0.01)  
 Max features: 25, num estimators: 10, accuracy: 0.32 (+/- 0.01)  
 Max features: 25, num estimators: 30, accuracy: 0.41 (+/- 0.01)  
 Max features: 25, num estimators: 50, accuracy: 0.43 (+/- 0.02)  
 Max features: 25, num estimators: 70, accuracy: 0.44 (+/- 0.02)  
 Max features: 25, num estimators: 90, accuracy: 0.44 (+/- 0.01)  
 Max features: 25, num estimators: 110, accuracy: 0.45 (+/- 0.01)  
 Max features: 25, num estimators: 130, accuracy: 0.45 (+/- 0.01)  
 Max features: 25, num estimators: 150, accuracy: 0.45 (+/- 0.01)  
 Max features: 25, num estimators: 170, accuracy: 0.46 (+/- 0.01)  
 Max features: 25, num estimators: 190, accuracy: 0.46 (+/- 0.00)  
 Max features: 30, num estimators: 10, accuracy: 0.33 (+/- 0.01)  
 Max features: 30, num estimators: 30, accuracy: 0.41 (+/- 0.01)  
 Max features: 30, num estimators: 50, accuracy: 0.42 (+/- 0.01)  
 Max features: 30, num estimators: 70, accuracy: 0.43 (+/- 0.01)  
 Max features: 30, num estimators: 90, accuracy: 0.44 (+/- 0.02)  
 Max features: 30, num estimators: 110, accuracy: 0.45 (+/- 0.01)  
 Max features: 30, num estimators: 130, accuracy: 0.45 (+/- 0.00)  
 Max features: 30, num estimators: 150, accuracy: 0.46 (+/- 0.01)  
 Max features: 30, num estimators: 170, accuracy: 0.46 (+/- 0.01)  
 Max features: 30, num estimators: 190, accuracy: 0.46 (+/- 0.01)  
 Max features: 35, num estimators: 10, accuracy: 0.33 (+/- 0.02)  
 Max features: 35, num estimators: 30, accuracy: 0.40 (+/- 0.01)



```


Max features: 35, num estimators: 50, accuracy: 0.43 (+/- 0.01)
Max features: 35, num estimators: 70, accuracy: 0.44 (+/- 0.01)
Max features: 35, num estimators: 90, accuracy: 0.44 (+/- 0.01)
Max features: 35, num estimators: 110, accuracy: 0.45 (+/- 0.01)
Max features: 35, num estimators: 130, accuracy: 0.45 (+/- 0.02)
Max features: 35, num estimators: 150, accuracy: 0.45 (+/- 0.01)
Max features: 35, num estimators: 170, accuracy: 0.45 (+/- 0.02)
Max features: 35, num estimators: 190, accuracy: 0.46 (+/- 0.01)
Max features: 40, num estimators: 10, accuracy: 0.33 (+/- 0.02)
Max features: 40, num estimators: 30, accuracy: 0.40 (+/- 0.01)
Max features: 40, num estimators: 50, accuracy: 0.43 (+/- 0.02)
Max features: 40, num estimators: 70, accuracy: 0.43 (+/- 0.02)
Max features: 40, num estimators: 90, accuracy: 0.44 (+/- 0.01)
Max features: 40, num estimators: 110, accuracy: 0.45 (+/- 0.01)
Max features: 40, num estimators: 130, accuracy: 0.45 (+/- 0.00)
Max features: 40, num estimators: 150, accuracy: 0.45 (+/- 0.01)
Max features: 40, num estimators: 170, accuracy: 0.46 (+/- 0.01)
Max features: 40, num estimators: 190, accuracy: 0.45 (+/- 0.01)
Max features: 45, num estimators: 10, accuracy: 0.34 (+/- 0.02)
Max features: 45, num estimators: 30, accuracy: 0.41 (+/- 0.02)
Max features: 45, num estimators: 50, accuracy: 0.42 (+/- 0.02)
Max features: 45, num estimators: 70, accuracy: 0.44 (+/- 0.02)
Max features: 45, num estimators: 90, accuracy: 0.44 (+/- 0.01)
Max features: 45, num estimators: 110, accuracy: 0.44 (+/- 0.01)
Max features: 45, num estimators: 130, accuracy: 0.45 (+/- 0.01)
Max features: 45, num estimators: 150, accuracy: 0.45 (+/- 0.02)
Max features: 45, num estimators: 170, accuracy: 0.45 (+/- 0.02)
Max features: 45, num estimators: 190, accuracy: 0.46 (+/- 0.01)

```

```

In [94]: import matplotlib.pyplot as plt
         from mpl_toolkits.mplot3d import Axes3D
         from matplotlib import cm
         fig = plt.figure()
         fig.clf()
         ax = fig.gca(projection='3d')
         x = rf_params[:,0]
         y = rf_params[:,1]
         z = rf_params[:,2]
         ax.scatter(x,y,z)
         ax.set_zlim(0.2, 0.5)
         ax.set_xlabel('Max features')
         ax.set_ylabel('Num estimators')
         ax.set_zlabel('Avg accuracy')
         plt.show()

```



In [ ]: